
0. Erlang

Sequential Programming

Facultat d'Informàtica de Barcelona (FIB)
Universitat Politècnica de Catalunya (UPC)

Why Erlang?

- We need a language to implement some of the algorithms that we will learn
- Erlang is a general-purpose language and runtime environment well suited for scalable distributed programming
 - Developed at Ericsson in late eighties
 - Functional Programming (e.g. Prolog, Lisp)
 - Available as open source
 - Built-in support for concurrency, distribution and fault tolerance

Why Erlang?

- Companies use Erlang in their production systems, e.g. Amazon, Yahoo!, Facebook, WhatsApp, T-Mobile, Motorola, Ericsson, ...
 - [https://en.wikipedia.org/w/index.php?title=Erlang_\(programming_language\)&oldid=754567048#Companies_using_Erlang](https://en.wikipedia.org/w/index.php?title=Erlang_(programming_language)&oldid=754567048#Companies_using_Erlang)
- Popular applications use Erlang, e.g. Chef, Ejabberd, CouchDB, GitHub, RabbitMQ, ...
 - [https://en.wikipedia.org/w/index.php?title=Erlang_\(programming_language\)&oldid=754567048#Software_projects_written_in_Erlang](https://en.wikipedia.org/w/index.php?title=Erlang_(programming_language)&oldid=754567048#Software_projects_written_in_Erlang)

Data Structures

- Literals
 - atoms: `foo`, `my_bar`, ...
 - numbers
 - integers: `10`, `-234`, `16#A42B`
 - floats: `17.3`, `-56.62`
 - `nil`: `[]`
 - `bool`: `true`, `false`

Data Structures

- Compound
 - tuples: used to store a **fixed** number of items
 - {123, def, abc}
 - {person, 'Joe', 'Armstrong'}
 - {abc, {def, 123}, jkl}
 - lists: used to store a **variable** number of items
 - []: empty list
 - [foo, 12, bar, zot]
 - "... " is short for the list of integers representing the ASCII codes of the enclosed within the quotes
 - "abcdefghi" is [97,98,99,100,101,102,103,104,105]

Variables

- Used to store values of data structures
- Procedure definition
 - No global scope
- Variables can only be bound **once**
 - Assigned a value when introduced
 - The value of a variable can never be changed
- Start with an upper case letter
 - Abc, A_var, Foo

Variables

- Assignment and Pattern matching
 - $A = 10$; Succeeds: binds A to 10
 - $B = \{z, \text{Foo}, 4\}$; Succeeds: binds B to $\{z, \text{Foo}, 4\}$
 - $\{B, C, D\} = \{10, \text{foo}, \text{bar}\}$
 - Succeeds: binds B to 10, C to foo and D to bar
 - $\{A, A, B\} = \{\text{abc}, \text{abc}, \text{foo}\}$
 - Succeeds: binds A to abc, B to foo
 - $\{A, A, B\} = \{\text{abc}, \text{def}, 123\}$; Fails
 - $[A, B, C] = [1, 2, 3]$
 - Succeeds: binds A to 1, B to 2, C to 3
 - $[A, B, C, D] = [1, 2, 3]$; Fails

Pattern matching

- Cons cell: [H | T]
- Used for pattern matching on lists
 - The pattern "[H|T] = List" extracts the head into "H" and tail into "T" of the list "List"
 - [A,B|C] = [1,2,3,4,5,6,7]
 - Succeeds: binds A = 1, B = 2, C = [3,4,5,6,7]
 - [H|T] = [1,2,3,4]
 - Succeeds: binds H = 1, T = [2,3,4]
 - [H|T] = [abc]
 - Succeeds: binds H = abc, T = []

Pattern matching

- $[H|T] = []$
 - Fails
- $\{A, _, [B|_], \{B\}\} = \{abc, 23, [22, x], \{22\}\}$
 - Succeeds: binds $A = abc$, $B = 22$
 - Note the use of " $_$ ", the anonymous (don't care) variable
 - It is used as a place holder where the syntax requires a variable, but the value of the variable is of no interest
 - If you know that X is bound to a tuple with three elements and you need to access the second, do:
 - $\{_, Y, _ \} = X$
 - Y is now a reference to the second element

Function calls

`module:func(Arg1, Arg2, ... Argn)`

- `Arg1 .. Argn` are any Erlang data structures
- The module/function name must be an atom

`math2:double(10).`

- Function calls can be nested

`math2:double(math2:double(2)).`

- A function can have zero arguments

`hello:world().`

Function calls

- Built In Functions (BIFs)
 - date()
 - time()
 - length([1,2,3,4,5])
 - size({a,b,c})
 - self()
 - ...
- Full list at <http://erlang.org/doc/index.html>

Function definition

func(Pattern1, Pattern2, ...) -> ... ;

func(Pattern1, Pattern2, ...) -> ... ;

...

func(Pattern1, Pattern2, ...) ->

- Clauses are scanned sequentially until a match is found
- When a match is found all variables occurring in the head become bound
- Variables are local to each clause

Function & module definition

-module(demo).

-export([double/1]).

double(X) -> times(X, 2).

times(X, N) -> X * N.

- Functions are defined within Modules and must be exported before they can be called from outside the module
 - double can be called from outside the module, times is local to the module
 - double/1 means the function double with one argument

Function examples

```
-module(mathStuff).  
-export([factorial/1, area/1]).
```

```
factorial(0) -> 1;
```

```
factorial(N) -> N * factorial(N-1).
```

```
area({square, Side}) -> Side * Side;
```

```
area({circle, Radius}) -> 3.14 * Radius * Radius;
```

```
area({triangle, A, B, C}) -> S = (A + B + C)/2,  
    math:sqrt(S*(S-A)*(S-B)*(S-C));
```

```
area(Other) -> {invalid_object, Other}.
```

Function examples: Conditionals

fac(N) ->

if

N == 0 -> 1;

N > 0 -> N*fac(N-1)

end.

sum([]) ->

0;

sum([H|T]) ->

H + sum(T).

sum(L) ->

case L of

[] ->

0;

[H|T] ->

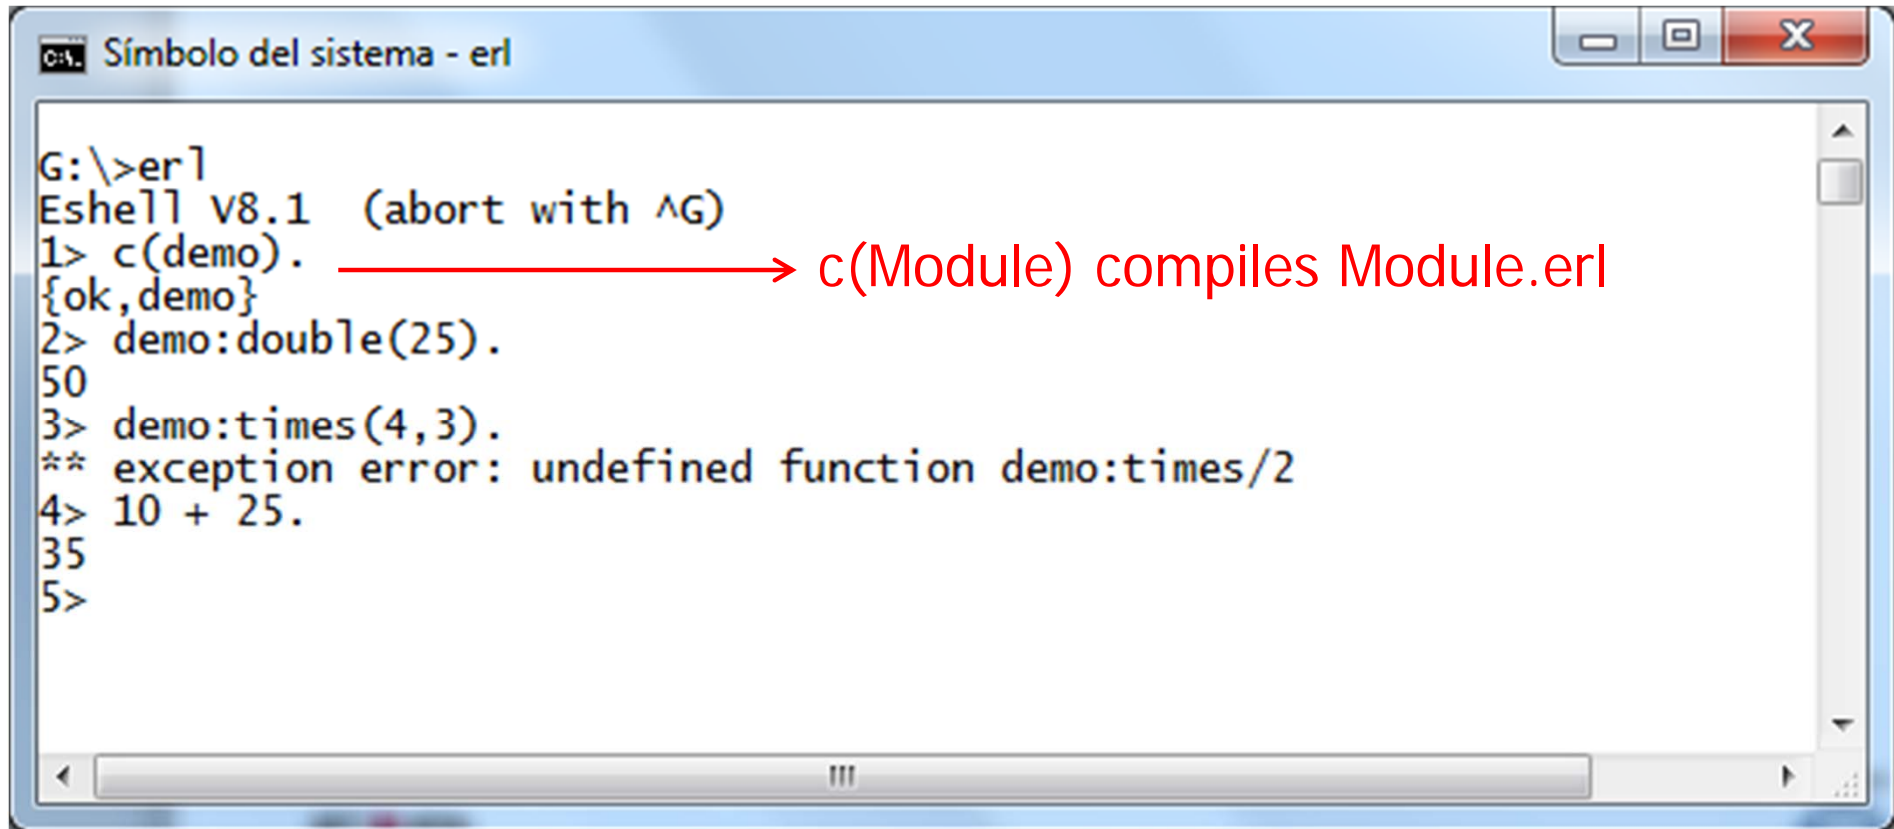
H + sum(T)

end.

Runtime system

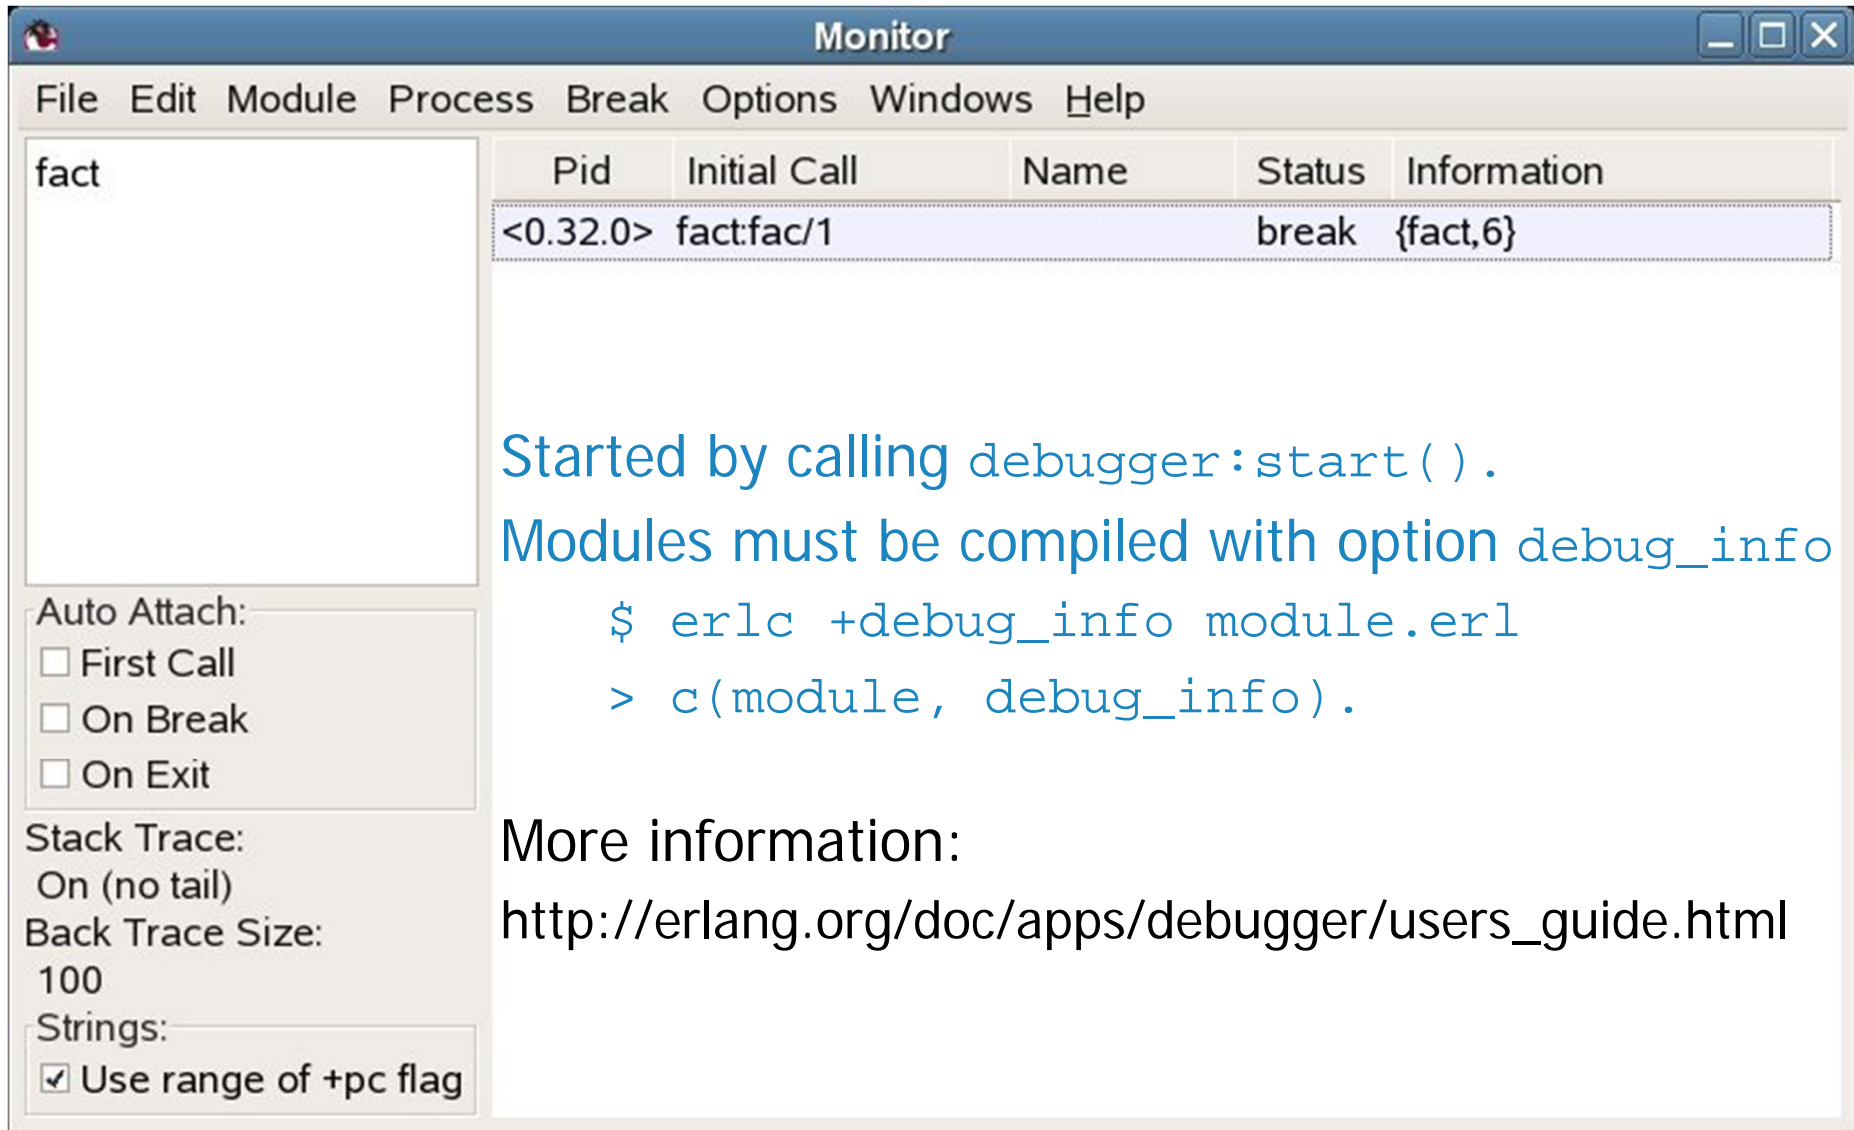
- The Erlang runtime system gives you an interactive shell
- In the shell you can compile and load modules and call functions
- Run it by itself, inside vim, emacs, or in a IDE such as Eclipse

Runtime system



```
G:\>erl
Eshell v8.1 (abort with ^G)
1> c(demo). —————→ c(Module) compiles Module.erl
{ok,demo}
2> demo:double(25).
50
3> demo:times(4,3).
** exception error: undefined function demo:times/2
4> 10 + 25.
35
5>
```

Debugger



The screenshot shows the Erlang Monitor window. The title bar is "Monitor". The menu bar includes File, Edit, Module, Process, Break, Options, Windows, and Help. The main area is divided into two panes. The left pane shows the module "fact". The right pane shows a table with columns: Pid, Initial Call, Name, Status, and Information. A single row is highlighted with a light blue background, showing the PID <0.32.0>, Initial Call fact:fac/1, Name, Status break, and Information {fact,6}. Below the table, there is a text area with the following text: "Started by calling `debugger:start()`. Modules must be compiled with option `debug_info`
`$ erlc +debug_info module.erl`
`> c(module, debug_info).`" Below the text area, there is a section titled "More information:" followed by the URL http://erlang.org/doc/apps/debugger/users_guide.html. On the left side of the window, there are several sections: "Auto Attach:" with three checkboxes (First Call, On Break, On Exit), all of which are unchecked; "Stack Trace:" with the text "On (no tail)"; "Back Trace Size:" with the value "100"; and "Strings:" with a checked checkbox "Use range of +pc flag".

| Pid | Initial Call | Name | Status | Information |
|----------|--------------|------|--------|-------------|
| <0.32.0> | fact:fac/1 | | break | {fact,6} |

fact

Auto Attach:

- ☐ First Call
- ☐ On Break
- ☐ On Exit

Stack Trace:
On (no tail)

Back Trace Size:
100

Strings:
☒ Use range of +pc flag

Started by calling `debugger:start()`.
Modules must be compiled with option `debug_info`
`$ erlc +debug_info module.erl`
`> c(module, debug_info).`

More information:
http://erlang.org/doc/apps/debugger/users_guide.html

0. Erlang

Concurrent Programming

Facultat d'Informàtica de Barcelona (FIB)
Universitat Politècnica de Catalunya (UPC)

Process creation

- Pid1 invokes:

Pid2 = spawn(Module, Function, [Args])

- Creates and starts the execution of a new process that runs Function
- Pid2 is process identifier of the new process
 - This is known only to process Pid1

Message passing

- To send a message you need the *process identifier* of the receiver

Pid ! Msg

- Sending a message is **asynchronous**
 - No acknowledgement
- self() - returns the Process Identity (Pid) of the process executing this function
- Process identifiers can be included in messages just like any data structure

P = spawn(wait, hello, []).

P ! "hello".

Message passing

- A process can suspend waiting for a message

receive

Msg -> Actions

end.

-module(wait).

-export([hello/0]).

hello() ->

receive

X -> io:format("message received: ~s~n", [X])

end.

Message passing

- A process will have an ordered sequence of received messages
 - All messages sent to a process are stored in its mailbox in the same order as they arrive

receive

Message1 -> Actions1 ;

Message2 -> Actions2

end.


- The first message that matches one of defined patterns will be delivered

Message passing

P1

P2

receive deliver



messages received
in FIFO order

asynchronous sending,
no acknowledgment

implicit deferral
of message
delivery

receive
m2 -> true
end,
receive
m1 -> true
end

Message passing

- ↑ One can select which messages to handle first
- ↓ Risk of forgetting messages that are left in a growing queue
- Receiving messages from a specific process
Pid ! {self(),abc}.

receive

{Pid,Msg} ->...

end.

Registered processes

- Register the process identifiers under names that are known to all processes

`P = spawn(wait, hello, []).`

`register(foo, P).`

`foo ! "hello".`


- Sending to a registered name is different to sending to a process id
 - Sending to a process id will always succeed
 - Even if the process is dead
 - A dead process will be de-registered and sending to a name without registered process will cause an exception

0. Erlang

Distributed Programming

Facultat d'Informàtica de Barcelona (FIB)
Universitat Politècnica de Catalunya (UPC)

Distributed programming



- Each Erlang instance is a different node in the distributed system

Node 1: 'gold@sdx.fib.upc.edu'

Node 2: 'silver@sdx.fib.upc.edu'

Distributed programming

- Create a process in a remote Erlang node

`P = spawn('gold@sdx.fib.upc.edu', M, F, [A]).`


`P ! Msg.`

- Connections are established automatically when another node is referenced
- Process IDs returned by spawn (or received in a message) can be used normally
 - Access transparent: local or remote are the same
 - Location transparent: you do not know where the process is located

Distributed programming

- Send a message to a process that is locally registered on a remote node

{Name, Node} ! Message.



{foo, 'gold@sdx.fib.upc.edu'} ! "hello".

Distributed programming

- We will use locally registered names, but Erlang offers also globally registered names
`global:register_name(foo, Pid).`
- Send a message to the process globally registered as `foo`
`global:send(foo, "Hello").`
- Connections must be established explicitly
`net_kernel:connect_node(Node).`

Distributed programming

- When you start Erlang you should make it network aware by providing a name

- Using long (-name) or short names (-sname) (*)

erl -name gold@sdx.fib.upc.edu / erl -name gold

- node(): gold@sdx.fib.upc.edu

erl -sname gold@sdx / erl -sname gold

- node(): gold@sdx

erl -name gold@127.0.0.1

(*) A node with a long name cannot communicate with a node with a short name

Distributed programming

- If someone connects to a node, it gets connected to all the other nodes
- Use cookies as a mechanism to differentiate clusters of nodes \Rightarrow Nodes with different cookies are not able to communicate together
 - setcookie mycookie
 - erlang:set_cookie(node(), mycookie)
 - Alternatively, you can have a file *.erlang.cookie* with the cookie in your home folder on all nodes

More information

- Erlang official website: <http://www.erlang.org/>
- 'An Erlang Primer' by Johan Montelius
 - <http://people.kth.se/~johanmon/dse/crash.pdf>
- 'Learn You Some Erlang for Great Good!'
 - <http://learnyousomeerlang.com/>
- 'Concurrent Programming in Erlang, Part I'
 - <http://erlang.org/download/erlang-book-part1.pdf>
- Elixir language: <http://elixir-lang.org/>
 - Runs on the Erlang runtime